Microsoft Fabric FAQ

James Serra
PRINCIPAL SOLUTION SPEC

Here are the FAQ's that I am hearing the most about Microsoft Fabric. Most of these answers were directly from various Fabric PMs:

I created a field-ready deck that hopefully will help you to learn about Fabric: Microsoft Fabric - James Serra.pptx. A presentation I gave on part of the deck is at Project Trident demo-20230427_160254-Meeting Recording.mp4. Note this deck has some NDA slides in it that can't be shown to a customer. I do have a customer-ready deck for a 60-minute presentation available here: Microsoft Fabric - James Serra - Customer.pptx. I highly recommend checking out the Fabric CAT Portal, and in particular the Fabric Mega-FAQ.

1) Will there be a demo built for us that we can use?

Sanjay Soni's team is building demo environments and material that will be available 6/15 at https://aka.ms/dreamdemos. This is difficult as we do not have APIs or Templates to quickly deploy artifacts (something can be done using Git integration). In the meantime, I recommend guiding customers to use https://learn.microsoft.com/en-us/fabric/get-started/end-to-end-tutorials

2) What do we tell customers who ask "how do we prevent certain users from using Fabric since it's all or nothing now"?

It is NOT an all or nothing. You can enable Fabric for a Tenant or for a specific capacity. Even after that, you can select specific security groups. See Enable Microsoft Fabric for your organization - Microsoft Fabric | Microsoft Learn

3) Who within MS are the people who should be talking to customers about Fabric - PBI specialists/CSAs or Data & AI specialists/CSAs?

Depends on where you are on the selling process. There will be guidance that will come out soon. This guidance will also include the Fabric CAT Team.

Data and AI specialist is the primary responsible specialist in the field for Fabric supported by DAI CSU and Intelligence TS.

4) What do I tell customers who want to build something now? Is it "Use Synapse and avoid using dedicated pools and do ADF instead of Synapse for pipelines and be aware Fabric does not support mapping data flows?"

Use Databricks today to start building your Delta lakehouse in an ADLS storage account. When Fabric becomes available, it will integrate natively with your existing Delta lakehouse.

Tell them they need to build the Lakehouse using parquet and delta lake, making sure data is stored in ADLS Gen 2 (so not conventional ADB managed tables). You can also tell them that, although Mapping Data Flows are not supported, the ones they build on Synapse will also be mounted.

In my opinion, the best ETL in Fabric is the Notebook based one. I hope that changes by GA time as we need a solid low-code ETL solution.

For me it's not that obvious... when you create a Lakehouse with Databricks and then shortcut it with Fabric then you will not benefit from V-Order which makes our Power BI Direct Lake so performant.

For customers who want to build something now (i.e. use GA tech) we can recommend the lake-first pattern as the easiest one for Fabric. You are right that vorder will not work, but the functionality will be there. Nothing in GA today produces vorder. At Fabric GA, they will be able to migrate their lake-first tooling to Fabric and benefit from vorder as well.

5) What do I tell users who see there is no Purview integration? How do they tell what is being created in Fabric?

This is a though one. When I get these questions, I tell them that first we are making sure we get a solid product out there (not a toy). I also tell them there is a group working on the "native" integration and you should expect to have a good one by GA time. Right now MIP integration lands first , which will be followed by other features like catalog / Data Loss Prevention extending to all fabric artifacts. This is currently work in progress

The DataHub is the built-in catalog for Fabric. It tracks everything that is created in Fabric, ensures MIPS etc. The upcoming Purview integration operates on top of the

existing DataHub.

Purview Hub in Fabric shows insights only about Fabric estate. OneLake Data Hub is only for Fabric assets. Purview catalog can include assets from other workloads outside of Fabric.

6) How can a customer that is using ADF today migrate that content to Fabric?

"Today" the response would be they "may to more than likely will" need to rebuild things within Fabric depending upon their current architecture and feature usage. In the future they will be able to "mount' existing ADF workloads into Fabric to utilize Fabric's compute model. If you have customer names where this is an ask please feel free to supply so we can leverage in engineering discussions and prioritization.

7) What should be the messaging for customers heavily invested in mapping data flows?

Power Query everything!!! - but on a more serious note Mapping Data Flows will be a part of ADF and they can "mount" these existing workloads into Fabric before considering refactoring. Any new projects they should look to what exists with the combination of Pipelines and Dataflows Gen2.

8) Do we have any internal hands-on lab / resources that I can leverage to build an entire flow?

End-to-end tutorials have been published on Docs: https://aka.ms/fabrice2etutorials

9) What is the difference between Lakehouse and <u>Warehouse</u>

In short, if you like Spark than use a Lakehouse, if you like <u>T-SQL</u> use Warehouse. The big difference is if built in Lakehouse then all updates must be done in Spark (can't do updates in T-SQL), or built in Warehouse then must do updates in T-SQL (can't do updates in Spark). But can read from both using Spark or T-SQL. Be aware of Delta lake shortcomings: no multi-table transaction support, limited T-SQL support, performance problem for trickle transactions.

A detailed table comparsion:
https://learn.microsoft.com/en-us/fabric/get-started/decision-guide-warehouse-lakehouse

10) Will we have the ability to limit access by individual products? For example, all certain users to create a Lakehouse but not a Warehouse?

No plans. For now it's all or nothing. However, you can do some limitations per workspace, but it does not go to the level of each product: [Roles in workspaces in Microsoft Fabric - Microsoft Fabric | Microsoft Learn](#)

11) Is it true that Fabric can be used for free until GA, and then free for 60 days after?

From https://learn.microsoft.com/en-us/fabric/get-started/fabric-trial "The Microsoft Fabric (Preview) trial includes access to the Fabric product experiences and the resources to create and host Fabric items. The Fabric (Preview) trial lasts until Fabric General Availability (GA), unless canceled. After GA, the Fabric (Preview) trial converts to the GA version and is extended for 60 days."
I believe it will continue to renew until we hit GA then the 60 days will be respected going forward. I don't know if that will autorenew or if they will need to go click a button until GA.
In today's session, Mihir confirmed that as long as customer use trial capacity at least once in a week, the trial period will keep on extending beyond 60 days and until GA.

12) Is there any report monitoring yet for a customer to determine what people are using Fabric?

Yes, see https://learn.microsoft.com/en-us/fabric/admin/feature-usage-adoption

13) Are there any details on Purview integration?

Yes, see https://learn.microsoft.com/en-us/fabric/governance/microsoft-purview-fabric and https://learn.microsoft.com/en-us/fabric/governance/use-microsoft-purview-hub. A video on Purview integration is at https://build.microsoft.com/en-US/sessions/b35ec76d-c6ee-4c71-ad28-7b3210a7f6be?source=sessions

15) Will we publicly list the roadmap so customers know what features are missing now and when they will be available?

Upcoming features will be announced via release notes.
A public roadmap will be released soon as the teams firm up details. Similar to the Power Platform Release Plan that exists today if they leverage that format

16) Can Databricks directly access OneLake?

Yes, see: https://learn.microsoft.com/en-us/fabric/onelake/onelake-azure-databricks

17) Is is possible to restrict how many Capacity Units can a workload and/or workspace consume?

We don't have these sorts of guard rails yet, but the product team is working on a set of "surge protection" features to address concerns like this.

18) Is there still MPP technology in Fabric?

Yes. It is using the POLARIS engine, a distributed query processing engine. More info on POLARIS is here

19) When will Fabric GA?

GA timing communication for Fabric:
- We will NOT publish the planned GA date publicly at Build or in the release plans
- **Under NDA**, we can tell customers we are working toward a GA date of Ignite (Nov 15-16) but that it will be solidified in the coming months based on feedback and usage of the public preview.
- In June, the engineering team will do a recalc of our semester plans based on our learnings from the public preview and execution velocity so far on our semester plans. This allows us space to add or remove scope and understand the feedback coming in.
- By July, we will lock on a GA date and scope that we are confident in. Based on this, we can decide if we want to go public with the date to allow customers to plan accordingly.

20) How will Fabric ensure that a client doesn't hit IOPS limitations when using multiple warehouse / spark and PBI engines against the same OneLake?

This will be handled by the OneLake service. From my understanding we will do some proactive monitoring of customer usage and raise the limits based on need. Some of this can be handled by the way that content/workloads are organized across workspaces - each workspace is a different storage account - compute is shared. It can also be addressed by ensuring that the delta tables are optimized (rather than a lot of smaller parquet segments). The CAT team will be working on

guidance and best practices and these are the sorts of nuanced things we are going to have to address.

21) Can I use SSMS?

Yes. Go to settings of Lakehouse or Warehouse and copy SQL connection string. In SSMS use "Azure Active Directory - Universal with MFA" for authentication. Once connected to SSMS, you will see all Lakehouses and Warehouses in your tenant

22) How do you choose the regions used in Fabric?

You define the regions when setting up Fabric capacities. Then you assign workspaces to the Fabric capacities. So, if you create a Fabric capacity called CAP1 is US West and create a workspace called WORK1 and specify it uses CAP1, then WORK1 will be in US West

23) How do capacities work now that there are two: Power BI Premium and Fabric Capacity?

All Power BI Premium capacities will continue to work as usual. If you want to use the new Fabric features, you will create a workspace that uses Fabric capacity. Eventually Power BI Premium capacities (P-SKUs) will be converted to Fabric capacities (the pricing structure is the same). The PBI v-cores will translate into Fabric CUs - eg. P1 = F64. Fabric capacities are available now: Announcing Microsoft Fabric capacities are available for purchase | Microsoft Fabric Blog | Microsoft Fabric

24) If a customer purchases a fabric capacity, will they start getting charged for using some of the new Fabric features, like Warehouse?

From Announcing Microsoft Fabric capacities are available for purchase | Microsoft Fabric Blog | Microsoft Fabric: "Until August 1, Fabric workloads other than Power BI will not count against your purchased capacity limits". If you purchase the capacity today through Azure the only workload that will count as consumption will be Power BI. That means you can consume 100% of the capacity with Power BI while still running Fabric workloads but starting on August 1st you're going to get charged when we start counting the consumption because you'll be using more than 100% of the capacity. You won't be charged for Fabric workloads on Fabric trials until GA. But if you associate a non-trial workspace with a Fabric capacity then you'll get charged starting August 1st

25) When will Fabric be available in the Azure government cloud?

It will be available 60-90 days after Fabric GA's in Commercial, provided all presence, capacities and dependencies are met. It will be FedRamp High when it GA's since its just a PBI upgrade and not a new product